

Part 2: Elementary Data Structures Implementation and Discussion

1. Implementation

Arrays (*full code available the file on Array.py*)

```
class MyArray:
    def __init__(self, capacity):
        self.capacity = capacity
        self.data = [None] * capacity # underlying storage
        self.length = 0 # number of actual stored elements

    def insert(self, index, value):
        if self.length == self.capacity:
            raise Exception("Array is full.")

        if index < 0 or index > self.length:
            raise Exception("Invalid index.")

        for i in range(self.length - 1, index - 1, -1):
            self.data[i + 1] = self.data[i]

        # Insert new element and increase Length
        self.data[index] = value
        self.length += 1

    def delete(self, index):
        if index < 0 or index >= self.length:
            raise Exception("Invalid index.")

        deleted_value = self.data[index]

        for i in range(index, self.length - 1):
            self.data[i] = self.data[i + 1]

        self.data[self.length - 1] = None
        self.length -= 1
        return deleted_value

    def access(self, index):
        if index < 0 or index >= self.length:
            raise Exception("Invalid index.")
        return self.data[index]

    def update(self, index, value):
        if index < 0 or index >= self.length:
            raise Exception("Invalid index.")
        self.data[index] = value

    def __str__(self):
        return str(self.data[:self.length])
```

```
● PS C:\Users\priyat> python arrays.py
===== ARRAY DEMO =====

Inserting 5 at index 0...
Array now: [5]

Inserting 10 at index 1...
Array now: [5, 10]

Inserting 7 at index 1 (between 5 and 10)...
Array now: [5, 7, 10]

Deleting value at index 1 (which is 7)...
Deleted value: 7
Array now: [5, 10]

Accessing element at index 1...
Value at index 1: 10

Updating index 1 to new value 20...
Array now: [5, 20]
PS C:\Users\priyat> █
```

Matrix (*full code available the file on Matrix.py*)

```
class MyMatrix:
    def __init__(self, rows, cols):
        # Create a rows x cols matrix initialized with zeros
        self.rows = rows
        self.cols = cols
        self.data = [[0 for _ in range(cols)] for _ in range(rows)]

    def access(self, r, c):
        self._validate(r, c)
        return self.data[r][c]

    def insert_row(self, r, row_values):
        if len(row_values) != self.cols:
            raise Exception("Row length mismatch.")
        self.data.insert(r, row_values)
        self.rows += 1

    def delete_row(self, r):
        self._validate(r, 0)
        del self.data[r]
        self.rows -= 1

    def insert_col(self, c, col_values):
        if len(col_values) != self.rows:
            raise Exception("Column length mismatch.")
        for i in range(self.rows):
            self.data[i].insert(c, col_values[i])
        self.cols += 1

    def delete_col(self, c):
        for i in range(self.rows):
            del self.data[i][c]
        self.cols -= 1

    def update(self, r, c, value):
        self._validate(r, c)
        self.data[r][c] = value

    def _validate(self, r, c):
        if r < 0 or r >= self.rows or c < 0 or c >= self.cols:
            raise Exception("Invalid matrix index.")
```

```
● PS C:\Users\priyat> python matrix.py
```

```
===== MATRIX DEMO =====
```

```
Initial 2x3 matrix:
```

```
[0, 0, 0]  
[0, 0, 0]
```

```
Updating (0,1) to 5...
```

```
[0, 5, 0]  
[0, 0, 0]
```

```
Updating (1,2) to 8...
```

```
[0, 5, 0]  
[0, 0, 8]
```

```
Inserting a new row [9, 9, 9] at index 1...
```

```
[0, 5, 0]  
[9, 9, 9]  
[0, 0, 8]
```

```
Deleting column 1...
```

```
[0, 0]  
[9, 9]  
[0, 8]
```

```
Accessing element at (1,1): 9
```

```
PS C:\Users\priyat> []
```

Stacks (*full code available the file on Stack.py*)

```
class Stack:
    def __init__(self):
        self.items = [] # Using Python List as the underlying array

    def push(self, item):
        """Add an item to the top of the stack."""
        self.items.append(item)

    def pop(self):
        """Remove and return the top item. Raise error if stack is empty."""
        if self.is_empty():
            raise IndexError("Pop from empty stack")
        return self.items.pop()

    def peek(self):
        """Return the top item without removing it."""
        if self.is_empty():
            raise IndexError("Peek from empty stack")
        return self.items[-1]

    def is_empty(self):
        """Check if the stack is empty."""
        return len(self.items) == 0

    def size(self):
        """Return the number of items in the stack."""
        return len(self.items)

# Example usage:
if __name__ == "__main__":
    stack = Stack()
    stack.push(10)
    stack.push(20)
    stack.push(30)
    stack.push(40)
    print("Top item:", stack.peek())
    print("Stack size:", stack.size())
    print("Popped item:", stack.pop())
    print("Is stack empty?", stack.is_empty())
```

```
PS C:\Users\priyat> python stack.py
❖ Top item: 40
Stack size: 4
Popped item: 40
Is stack empty? False
PS C:\Users\priyat> []
```

Queues (*full code available the file on Queue.py*)

```
class Queue:
    def __init__(self):
        self.items = []
    def enqueue(self, item):
        """Add an item to the rear of the queue."""
        self.items.append(item)

    def dequeue(self):
        """Remove and return the front item. Raise error if queue is empty."""
        if self.is_empty():
            raise IndexError("Dequeue from empty queue")
        return self.items.pop(0) # Remove the first element

    def front(self):
        """Return the front item without removing it."""
        if self.is_empty():
            raise IndexError("Front from empty queue")
        return self.items[0]

    def is_empty(self):
        """Check if the queue is empty."""
        return len(self.items) == 0

    def size(self):
        """Return the number of items in the queue."""
        return len(self.items)

# Example usage:
if __name__ == "__main__":
    queue = Queue()
    queue.enqueue(10)
    queue.enqueue(20)
    queue.enqueue(30)
    queue.enqueue(40)
    print("The queue is:", queue.items)
    print("Front item:", queue.front())
    print("Queue size:", queue.size())
    print("Dequeued item:", queue.dequeue())
    print("Is queue empty?", queue.is_empty())
```

```
● PS C:\Users\priyat> python queue.py
The queue is: [10, 20, 30, 40]
Front item: 10
Queue size: 4
Dequeued item: 10
Is queue empty? False
❖ PS C:\Users\priyat> █
```

Linked List (*full code available the file on [Linked_List.py](#)*)

```
class LinkedList:
    def __init__(self):
        self.head = None # Start of the linked list

    def insert_at_head(self, data):
        """Insert a new node at the beginning of the list."""
        new_node = Node(data)
        new_node.next = self.head
        self.head = new_node

    def insert_at_tail(self, data):
        """Insert a new node at the end of the list."""
        new_node = Node(data)
        if self.head is None:
            self.head = new_node
            return
        current = self.head
        while current.next:
            current = current.next
        current.next = new_node

    def delete(self, data):
        """Delete the first node with the given data."""
        current = self.head
        prev = None
        while current:
            if current.data == data:
                if prev:
                    prev.next = current.next
                else:
                    self.head = current.next
                return True # Node deleted
            prev = current
            current = current.next
        return False # Node not found

    def search(self, data):
        """Search for a node with the given data."""
        current = self.head
        while current:
            if current.data == data:
                return True
            current = current.next
        return False
```

```
● PS C:\Users\priyat> python linkedList.py
10 -> 20 -> 30 -> 40
Search 20: True
Delete 20: True
10 -> 30 -> 40
✿ PS C:\Users\priyat> █
```

Rooted Trees (*full code available on [Rooted_Trees.py](#)*)

```
class RootedTree:
    def __init__(self, root_data):
        self.root = TreeNode(root_data)

    def _search(self, node, value):
        if node.data == value:
            return node
        for child in node.children:
            found = self._search(child, value)
            if found:
                return found
        return None

    # Access
    def access(self, value):
        """Return the node with the given value, or None."""
        return self._search(self.root, value)

    # Insertion
    def insert(self, parent_value, new_value):
        """Insert a new node as a child of a node with parent_value."""
        parent_node = self._search(self.root, parent_value)
        if parent_node is None:
            raise ValueError(f"Parent '{parent_value}' not found.")

        parent_node.children.append(TreeNode(new_value))

    # Update
    def update(self, old_value, new_value):
        """Update the data of an existing node."""
        node = self._search(self.root, old_value)
        if node is None:
            raise ValueError(f"Node '{old_value}' not found.")

        node.data = new_value

    # Delete
    def delete(self, value):
        if self.root.data == value:
            raise ValueError("Cannot delete the root node.")

        # BFS to find the parent
        queue = [self.root]
```

```
● PS C:\Users\priyat> python Rooted_Tree.py
Tree Structure:
A
B
  D
  E
C
  F

DFS Traversal:
A
B
D
  B
    D
    E
C
  F
```

2. Performance Analysis

Arrays

Arrays are one of the most fundamental data structures in computer science. They store elements in contiguous memory locations, allowing efficient random access. This following examines the time complexity of four primary array operations: access, update, insertion, and deletion.

1. Access Operation

Accessing an element in an array using its index is extremely efficient.

- **Time Complexity:** $O(1)$
- **Explanation:** Since array elements are stored in contiguous memory, the address of any element can be calculated directly using its index. This makes the access operation constant time regardless of the array size.

2. Insertion Operation

The time complexity of inserting an element in the array depends on the position of insertion.

- **Time Complexity:** $O(1)$

- **Insert at the end :**
 - Static array: $O(1)$ if space is available; otherwise, insertion is not possible.
 - Dynamic array: $O(1)$ amortized; occasionally $O(n)$ if resizing is required.
- **Insert at the Beginning or Middle - $O(n)$:** When inserting at any position other than the end, all subsequent elements must be shifted by one position to make space. This shifting requires time proportional to the array size.

3. Deletion Operation

Like insertion, the time for deletion also depends on the location of the element being removed.

- Delete the last element: $O(1)$: No shifting is required, so deletion is constant time.
- Delete from the Beginning or Middle $O(n)$: all elements that come after the removed element must be shifted left to fill the gap, resulting in linear time complexity.

4. Updating Operation

Updating an element in an array using its index is extremely efficient.

- **Time Complexity:** $O(1)$
- **Explanation:** Just like access, updating a value involves computing the address from the index and writing the new value, which takes constant time.

Arrays provide fast access and update operations due to their contiguous memory layout.

However, insertion and deletion can be inefficient when they occur at positions other than the end of the array because of the required element shifting. Understanding these complexities helps in selecting the appropriate data structure based on the operations required by a specific application.

Matrices

A matrix is a two-dimensional data structure consisting of rows and columns. Like arrays, matrices store elements in contiguous memory (row-major or column-major order), allowing

efficient access and updates. This report summarizes the time complexity of common matrix operations, assuming an $m \times n$ matrix.

➤ Access Operation

- **Time Complexity:** $O(1)$
- **Explanation:** Accessing an element at row i and column j requires computing its memory address using the formula (row-major or column-major indexing). This calculation is constant-time.

➤ Insertion Operation

Insertion in a matrix depends on whether we are inserting a row or column:

- **Insert a row/column: $O(m \times n)$:**
 - In most implementations, matrices are fixed-size (static), so inserting a new row or column requires allocating a new matrix and copying all elements..
 - In dynamic or resizable matrix structures, insertion may be slightly faster for rows or columns at the end, but element copying is still generally required.
- **Insert an element at a specific position: $O(1)$**
 - If the matrix size is fixed, replacing a value is constant-time.
 - If the matrix needs to grow dynamically, copying may be required, making it $O(m \times n)$ in the worst case.

➤ Deletion Operation

Deletion also depends on rows, columns, or individual elements:

- **Delete a row/column: $O(m \times n)$**
 - Requires shifting or copying elements to maintain the matrix structure.
- **Delete an element: $O(1)$**
 - Setting a single element to a null or zero value is constant-time.

➤ Updating Operation

- **Time Complexity:** O(1)

- **Explanation:** Updating an element at a given position involves writing to its memory address, which is constant time.

Matrix operations such as access and update are efficient due to direct indexing. However, inserting or deleting entire rows or columns is costly because it often involves creating a new matrix and copying elements. Understanding these complexities is crucial when designing algorithms that manipulate large matrices.

Stacks

A stack is a linear data structure that follows the Last-In-First-Out (LIFO) principle. Elements are inserted and removed from the top of the stack only. Stacks can be implemented using arrays or linked lists, and this affects performance for some operations.

➤ Push Operation (Insert at Top)

- **Time Complexity:** O(1)
- **Explanation:** Adding an element to the top of the stack is a constant-time operation, whether implemented with an array (if space is available) or a linked list.
- **Array-based stack note:** If the array is dynamic and resizing is needed, the worst-case complexity becomes O(n), but amortized complexity remains O(1).

➤ Pop Operation (Removal from top)

- Time Complexity: O(1)
- **Explanation:** Removing the top element is constant-time because no shifting is required.

➤ Peek Operation (Access Top Element)

- **Time Complexity:** O(1)

- **Explanation:** Retrieving the top element without removing it only requires accessing a single memory location.
- IsEmpty / IsFull Operations
- **Time Complexity:** O(1)
 - **Explanation:** Checking whether the stack is empty or full (in the case of array-based stacks) is a simple comparison operation.

Stack operations are highly efficient because they are restricted to one end of the data structure (the top). Both array-based and linked-list implementations support O(1) time for push, pop, and peek operations, though dynamic arrays may occasionally require O(n) resizing. Stacks are ideal for algorithms requiring LIFO behavior, such as expression evaluation, recursion, and backtracking.

Queue

A queue is a linear data structure that follows the First-In-First-Out (FIFO) principle. Elements are inserted at the rear and removed from the front. Queues may be implemented using arrays, circular arrays, or linked lists.

- **Enqueue Operation (Insert from rear)**
- Time Complexity O(1)
 - **Explanation:** In a properly implemented queue such as one using a circular array or a linked list, adding an element to the rear requires only updating one or two pointers. No shifting is required, making the operation constant-time. If a simple array without circular indexing is used, inserting at the rear may eventually require shifting elements when the front moves forward, resulting in O(n) time. This is why circular arrays are preferred.
- **Dequeue Operation (remove from front)**
- Time Complexity O(1)
 - **Explanation:** Removing an element from the front simply involves moving the front pointer to the next element. Since no elements need to be shifted, the operation remains constant-time for both array-based and linked-list queues.

➤ Front/Peak Operation

- Time Complexity O(1)
- **Explanation:** Accessing the front element only requires referencing the index or pointer that represents the front of the queue.

➤ IsEmpty / IsFull

- Time Complexity O(1)
- These are simple checks that compare pointer values or counters.

Linked List

A linked list is a dynamic data structure composed of nodes, where each node stores data and a pointer to the next node. Unlike arrays, linked lists do not store elements in contiguous memory, enabling efficient insertions and deletions but slower element access.

➤ Access Operation

- Time Complexity: O(n)
- **Explanation:**
Linked lists do not support random indexing. To access the k -th element, the list must be traversed from the head node until the target node is reached.

➤ Insertion Operation

- **Insert at beginning (head): O(1)** Requires creating a new node and adjusting the head pointer.
- **Insert at end:**
 - O(1) if a tail pointer exists
 - O(n) without a tail pointer, because traversal to the last node is needed
- **Insert at a specific position: O(n)** Requires traversing the list to find the insertion point. The pointer changes themselves are constant-time.

➤ Update Operation

- Time Complexity: $O(n)$
- **Explanation:** Updating a node requires first locating it, which takes $O(n)$. The actual update is constant-time, but the traversal dominates the cost.

➤ Deletion Operation

- Delete at beginning: $O(1)$ - Only the head pointer needs to be updated.
- Delete at end: $O(n)$ unless a pointer to the second-last node is maintained
Searching for the node before the last requires full traversal.
- Delete at a specific position: $O(n)$ - The list must be traversed to find the node to remove and its predecessor.

Rooted Tree

A rooted tree is a hierarchical data structure consisting of nodes connected by edges, where one node is designated as the root. Each node may have zero or more child nodes, forming parent–child relationships. Rooted trees are widely used in file systems, expression parsing, hierarchical modeling, and search algorithms.

Because tree nodes are not stored in contiguous memory and typically use pointers, the time complexity of operations depends heavily on the structure of the tree (balanced vs. unbalanced) and on whether parent pointers are maintained.

➤ Accessing a Node

- **Time Complexity:** $O(h)$, where h is the height of the tree
- **Explanation:**

To reach a node, traversal must begin from the root and follow pointers down through its children. The height determines how many steps are needed.

- For balanced trees, $h = O(\log n)$
- For skewed/unbalanced trees, $h = O(n)$

➤ Searching a Node

- **Time Complexity:** $O(n)$ in general trees

- **Explanation:**

In a general rooted tree without ordering (e.g., no binary search property), every node may need to be visited to locate a specific value.

- **Best case:** Found at or near the root → $O(1)$
- **Worst case:** Value does not present or in deepest node → $O(n)$

➤ **Insertion Operation**

- **Time Complexity:** $O(1)$ to insert **after locating the parent**, otherwise $O(n)$
- **Explanation:** Inserting a node into a rooted tree is done by:

- Finding the parent node
- Adding the new node as one of its children

The actual linking step (pointer assignment) is constant-time.

However, if the parent is not already known, locating it requires traversal
→ $O(n)$.

➤ **Deletion Operation**

- **Time Complexity:** $O(1)$ to detach a known node; $O(n)$ if searching is required
- **Explanation:**

- Finding the node ($O(n)$ if search is needed)
- Adjusting pointers ($O(1)$)

In many tree types, deleting a node also requires handling its children:

- **Delete entire subtree:** $O(\text{size of subtree})$
- **Promote or reattach children:** $O(k)$, where k is number of children

In general rooted trees, deletion cost varies based on strategy.

➤ **Traversal Operation**

Common traversals include preorder, inorder, postorder, and level-order.

- **Time Complexity:** $O(n)$
- Every traversal algorithm must visit all nodes exactly once.

The trade-offs between using arrays versus linked lists for implementing stacks and queues.

When implementing stacks and queues, arrays and linked lists are the two most common underlying data structures. Both support efficient insertion and deletion at one end, but they differ significantly in memory usage, performance behavior, and flexibility. The main trade-offs are outlined below.

1. Arrays for Stacks and Queues

Stacks

- **Push and Pop:** $O(1)$ on average.
- **Dynamic resizing:** Occasional $O(n)$ when the array needs to expand.
- **Advantages:** Fast operations due to contiguous memory, low per-element memory overhead, and simple implementation.
- **Disadvantages:** Fixed-size arrays risk overflow; resizing introduces unpredictability.

Queues

- **Enqueue and Dequeue:** $O(1)$ if implemented as a **circular buffer**.
- **Shifting elements:** If a simple array is used, dequeue may require $O(n)$ time due to shifting.
- **Advantages:** Cache-friendly, predictable when capacity is sufficient, and simple to implement.
- **Disadvantages:** Resizing or shifting elements can degrade performance; memory may be wasted if preallocated space is not fully used.

Summary: Arrays work well when the maximum size is predictable, and operations are mostly at one end (stack) or managed with circular logic (queue). Performance is excellent in practice due to contiguous storage, but dynamic resizing and fixed capacity can be limitations.

2. Linked Lists for Stacks and Queues

Stacks

- **Push and Pop:** Always O(1), performed at the head of the list.
- **Advantages:** Flexible size, no resizing needed, and predictable operation time.
- **Disadvantages:** Extra memory per node for pointers; slower due to pointer dereferencing and poor cache locality.

Queues

- **Enqueue at tail, Dequeue at head:** O(1) with head and tail pointers.
- **Advantages:** Grows dynamically without resizing, consistent O(1) performance, memory used only for actual elements.
- **Disadvantages:** Extra memory for pointers, scattered nodes reduce cache performance, and memory allocation/deallocation can add overhead.

Summary: Linked lists are ideal for stacks and queues when growth is unpredictable or maximum size is unknown. They ensure predictable O(1) push, pop, enqueue, and dequeue, but incur extra memory overhead and slower practical performance.

Trade-offs Comparison

Feature	Arrays	Linked Lists
Stack Operations	Push/Pop O(1) amortized, resizing occasional O(n)	Push/Pop O(1) always
Queue Operations	Circular buffer: O(1); simple array: O(n) for dequeue	Enqueue/Dequeue O(1) with head/tail pointers
Memory Usage	Low; may waste capacity when reallocated	Higher per element (pointers)
Cache Performance	Excellent (contiguous memory)	Poor (nodes scattered)

Flexibility	Limited by fixed size or costly resizing	Dynamic, grows/shrinks as needed
Predictability	Resizing spikes may occur	Consistent O(1) for all operations

Conclusion

- **Use arrays** for stacks and queues when:
 - Maximum size is predictable.
 - High speed and memory efficiency are required.
 - Resizing cost or fixed capacity is acceptable.
- **Use linked lists** for stacks and queues when:
 - Size is dynamic and unpredictable.
 - Constant-time insertion/removal is critical.
 - Memory overhead is not a major concern.

Overall: Arrays offer better performance and cache utilization, whereas linked lists provide flexible growth and consistent O(1) operations, making them suitable for different use cases in stacks and queues.

Compare the efficiency of different data structures in specific scenarios.

Data structures differ in their efficiency depending on the operations required and the context in which they are used. Choosing the right data structure is critical for performance, memory usage, and scalability. Below is a detailed comparison of common data structures in specific scenarios.

1. Arrays

Best suited for:

- **Random access:** Arrays provide **O(1)** access by index, making them ideal when frequent retrieval of elements at known positions is required.
- **Static datasets:** When the number of elements is fixed or rarely changes, arrays minimize memory overhead and maximize cache efficiency.

Limitations:

- **Insertion/deletion in middle:** $O(n)$, because elements must be shifted.
- **Dynamic resizing:** Dynamic arrays may incur **$O(n)$ resizing cost**, though amortized insertion remains $O(1)$.

Example scenario: Lookup tables, buffers, and static lists where frequent indexed access is required.

2. Linked Lists

Best suited for:

- **Frequent insertions and deletions:** $O(1)$ at the head or tail (for singly/doubly linked lists with pointers).
- **Unpredictable size:** Memory grows dynamically without preallocating a fixed size.

Limitations:

- **Access by index:** $O(n)$, as traversal is required.
- **Cache performance:** Poor, because nodes are scattered in memory.

Example scenario: Implementing dynamic stacks and queues where growth is unpredictable, or adjacency lists in graphs.

3. Stacks and Queues

These are **abstract data structures** often implemented using arrays or linked lists.

Stacks:

- Efficient when implemented using arrays or linked lists.
 - Array: $O(1)$ amortized push/pop, excellent cache locality.
 - Linked list: $O(1)$ push/pop, flexible size, predictable operation.

Queues:

- Array-based circular buffers: $O(1)$ enqueue/dequeue, memory-efficient.
- Linked list: $O(1)$ enqueue/dequeue, no resizing needed, suitable for unbounded or dynamic queues.

Example scenario:

- Stack: Undo functionality, expression evaluation.
- Queue: Task scheduling, breadth-first search.

4. Hash Tables

Best suited for:

- **Fast search, insert, delete by key:** $O(1)$ average-case time complexity.
- **Dynamic datasets with unpredictable size:** Handles sparse or large key spaces efficiently.

Limitations:

- **Memory overhead:** Needs extra space for hash table and handling collisions.
- **Worst-case lookup:** $O(n)$ if many collisions occur (rare with good hash function).

Example scenario: Dictionaries, symbol tables, caching, and lookup-heavy applications.

5. Trees

Binary Search Tree (BST)

- **Best for:** Ordered data, fast search, insertion, and deletion — $O(\log n)$ on average.
- **Limitations:** Unbalanced BSTs may degrade to $O(n)$.

AVL / Red-Black Tree

- Self-balancing trees ensure $O(\log n)$ for search, insert, and delete in all cases.
- More complex to implement, but guaranteed efficiency.

Example scenario: Implementing priority-based structures, sorted datasets, and range queries.

In conclusion, the efficiency of a data structure depends on the specific operations and context in which it is used. Arrays are highly efficient for static datasets requiring frequent random access, while linked lists excel in dynamic scenarios with frequent insertions and deletions. Stacks and queues can be implemented efficiently using either arrays or linked lists, depending on whether predictable growth or cache-friendly performance is prioritized. Hash tables offer fast key-based access, trees maintain ordered data with predictable search and update times, and graph representations must be chosen based on density and traversal needs. Ultimately, selecting the most suitable data structure requires balancing factors such as operation frequency, memory usage, performance predictability, and the specific requirements of the application.

3. Discussion

Provide a discussion on the practical applications of these data structures in real-world scenarios.

1. Arrays

Description and Characteristics:

Arrays store elements in contiguous memory locations, which allows constant-time ($O(1)$) access by index. They are most efficient for scenarios where frequent random access is needed and the size of the dataset is either known or grows infrequently.

Applications:

- Lookup tables: Arrays provide fast $O(1)$ access by index, making them ideal for mapping values in graphics, sound processing, or mathematical computation.
- Buffers: Arrays are used for audio, video, and network buffers, where sequential read/write operations are frequent.
- Matrices: Arrays efficiently represent 2D or multi-dimensional data for scientific computing, spreadsheets, and grid-based systems.

Real-World Examples:

- Graphics processing: Pixel data for images, sprites in games, and frame buffers in video rendering.
- Numerical simulations: Scientific computations, weather modeling, or physics simulations using multidimensional arrays.
- Gaming: Tile-based maps in games, collision detection grids, and pathfinding matrices.

2. Linked Lists

Description and Characteristics:

Linked lists store elements in nodes connected via pointers. Unlike arrays, they can **grow or shrink dynamically** without requiring contiguous memory, and insertion/deletion at the head or tail is **$O(1)$** .

Applications:

- Dynamic memory management: Useful when the number of elements is unknown or highly variable.
- Undo/redo stacks in software: Each action is stored in a linked list for efficient insertion and removal.
- Graph representation: Sparse graphs are represented efficiently using adjacency lists.

Real-World Examples:

- Text editors: Storing actions for undo/redo functionality.
- Task scheduling systems: Jobs in operating systems or server queues with unpredictable lengths.
- Sparse network modeling: Adjacency lists for social networks or road maps.

3. Stacks

Description and Characteristics:

Stacks follow a Last-In-First-Out (LIFO) order. They are often implemented using arrays (for fixed or dynamic size) or linked lists (for dynamic size).

Applications:

- Function call management: Programming languages use stacks to manage function calls and return addresses.
- Expression evaluation: Compilers and calculators use stacks to evaluate arithmetic expressions in postfix/prefix notation.
- Undo mechanisms: Track user actions for undo/redo in applications.

Real-World Examples:

- Web browsers: Back and forward navigation uses stacks.
- Compilers: Parsing nested expressions and managing recursion.
- Text editing software: Undo/redo functionality in Microsoft Word or Photoshop.

4. Queues

Description and Characteristics:

Queues follow a First-In-First-Out (FIFO) order. They are implemented using arrays (with circular buffers) or linked lists (dynamic queues).

Applications:

- Task scheduling: Operating systems use queues for managing processes and threads.
- Resource management: Queues handle print jobs, network packets, and CPU jobs.
- Breadth-first traversal (BFS): Essential in graph traversal algorithms.

Real-World Examples:

- Network routers: Packet queues manage data transmission efficiently.
- Printer servers: Print jobs are processed in arrival order.
- Customer service systems: Ticketing systems or call centers process requests in FIFO order.

5. Trees

Binary Search Trees (BST):

- **Applications:** Efficient search, insertion, and deletion of ordered data.
- **Examples:** Contact lists, word dictionaries, or event logs.

Balanced Trees (AVL, Red-Black):

- **Applications:** Guarantee $O(\log n)$ operations for large dynamic datasets.
- **Examples:** Database indices, file systems, memory management systems.

Heaps:

- **Applications:** Implement priority queues and efficiently extract the minimum/maximum element.
- **Examples:** CPU scheduling, shortest path algorithms (Dijkstra), and data compression (Huffman coding).

6. Graphs

Description and Characteristics:

Graphs model networks of interconnected nodes. They can be represented as adjacency lists (sparse graphs) or adjacency matrices (dense graphs).

Applications:

- Social networks: Represent users and connections.
- Navigation systems: Road networks and GPS for shortest path computation.
- Dependency management: Represent relationships in tasks, software packages, or projects.

Real-World Examples:

- Facebook/LinkedIn: Users as nodes, friendships as edges.
- Google Maps/Waze: Road intersections as nodes, roads as edges.
- Project management tools: Task dependencies in Gantt charts or workflow systems

Data Structure	Key Applications	Real-World Examples
Array	Random access, buffers, matrices	Graphics pixels, numerical simulations, game maps
Linked List	Dynamic growth, frequent insert/delete	Undo/redo in editors, adjacency lists, task queues
Stack	LIFO operations	Call stack, browser history, expression evaluation
Queue	FIFO operations	Print servers, task scheduling, network packet queues
BST / Balanced Trees	Ordered data	Contact lists, database indices, event logs

Different data structures excel in different real-world scenarios. Arrays are ideal for fixed-size or lookup-intensive datasets due to fast random access. Linked lists handle dynamic datasets with frequent insertions and deletions. Stacks and queues manage ordered operations, supporting tasks like function calls, expression evaluation, and scheduling. Hash tables provide near-instant key-based retrieval, essential in caching and database systems. Trees maintain hierarchical or sorted data, and heaps optimize priority-based operations. Graphs model complex networks and relationships, essential in social media, navigation, and dependency analysis. Choosing the right data structure depends on operation frequency, memory constraints, and performance requirements, ensuring efficient, scalable, and reliable systems.

Highlight scenarios where one data structure may be preferred over another due to factors like memory usage, speed, and ease of implementation.

The choice of a data structure depends heavily on specific requirements of the application, such as how often elements are inserted or deleted, how quickly data needs to be accessed, and how much memory is available. Below is a scenario-based discussion comparing common data structures.

1. Arrays vs. Linked Lists

Factor	When Arrays are Preferred	When Linked Lists are Preferred
Memory Usage	Efficient for small elements and fixed-size datasets; contiguous storage reduces overhead	More memory-intensive due to pointers per node but ideal for datasets with highly variable sizes
Speed	Fast for random access ($O(1)$) and iteration due to cache locality	Slower access ($O(n)$) due to pointer traversal, but insert/delete at head/tail is $O(1)$

Ease of Implementation	Simple to implement; resizing for dynamic arrays adds minor complexity	Slightly more complex due to node and pointer management
Scenario Examples	Storing pixel data, static lookup tables, audio buffers	Dynamic stacks/queues, undo/redo buffers, adjacency lists for sparse graphs

Summary: Arrays are preferred when fast access and low memory overhead are critical, while linked lists excel when frequent insertions/deletions and dynamic sizing are required.

2. Stacks vs. Queues

Factor	Stack	Queue
Memory Usage	Minimal if implemented with arrays; slightly more with linked lists	Minimal if using circular buffers; slightly more with linked lists
Speed	O(1) push/pop with arrays or linked lists	O(1) enqueue/dequeue with circular arrays or linked lists
Ease of Implementation	Very simple with arrays; moderate with linked lists	Circular arrays require slightly more logic; linked lists are straightforward
Scenario Examples	Function call management, browser history, expression evaluation	Task scheduling, print servers, network packet management

Summary: Both can be implemented efficiently, but stacks are ideal for LIFO tasks, while queues suit FIFO processing, with implementation choices influenced by whether the size is fixed or dynamic.

3. Hash Tables vs. Trees

Factor	Hash Table	Tree (BST / Balanced Tree)
Memory Usage	Extra memory for hash buckets and collision handling	Pointer overhead per node; balanced trees may require additional storage for balance info
Speed	O(1) average-case for search, insert, delete	O(log n) for balanced trees; O(n) worst-case for unbalanced BST
Ease of Implementation	Moderate (hash function design, collision handling)	Moderate to complex, especially for self-balancing trees
Scenario Examples	Fast key-value lookups (caching, symbol tables, dictionaries)	Ordered data, range queries, database indexing, event logs

Summary: Use hash tables when fast key-based access is the priority and ordering is not required. Use trees when data must remain ordered or range queries are needed.

4. Graphs – Adjacency Matrix vs. Adjacency List

Factor	Adjacency Matrix	Adjacency List
Memory Usage	$O(n^2)$, expensive for sparse graphs	$O(n + e)$, memory-efficient for sparse graphs
Speed	$O(1)$ edge lookup	$O(k)$ edge lookup ($k = \text{degree of node}$)
Ease of Implementation	Simple; easy to check if an edge exists	Slightly more complex; traversing neighbors requires iteration

Scenario Examples	Dense networks, computer simulations, connectivity matrices	Social networks, road maps, dependency graphs
-------------------	---	---

Summary: Use adjacency matrices for dense graphs where fast edge lookup is important, and adjacency lists for sparse graphs to save memory.

In practice, selecting a data structure requires balancing memory usage, operational speed, and implementation complexity. Arrays are ideal for fast random access and memory efficiency, whereas linked lists support dynamic growth and frequent insertions/deletions. Stacks and queues efficiently manage LIFO and FIFO tasks, with implementation choice depending on size predictability. Hash tables excel in rapid key-value retrieval, while trees maintain ordered datasets for searches and range queries. Graph representations must be chosen based on density and traversal requirements, and heaps are optimal for priority-based operations. Ultimately, the choice of data structure should align with the specific operational demands, dataset characteristics, and system constraints of the application.